

DHYEY RAJESH SHAH

Dallas, TX · (945) 241-6194 · dhyeyshahr@gmail.com · [linkedin.com/in/dhyeyshah28801](https://www.linkedin.com/in/dhyeyshah28801) · github.com/dhyeyshah28801 · dhyeyshah28801.github.io/portfolio

Software engineer with 3+ years building backend platforms, AI/LLM systems, and cloud infrastructure across fintech, e-commerce, and supply-chain — 60% faster APIs, 80% fewer deployment failures, 99.9% availability. M.S. Computer Science (3.93 GPA).

EDUCATION

M.S. Computer Science, University of Texas at Dallas — GPA 3.93/4.0

Aug 2024 – May 2026

B.E. Computer Engineering, Gujarat Technological University — GPA 9.5/10

Jun 2019 – Jun 2023

TECHNICAL SKILLS

Languages: Java, Python, TypeScript, JavaScript, SQL, C/C++, Dart

Backend & Cloud: Spring Boot, Node.js, Express, REST APIs, Microservices, AWS (Lambda, S3, CloudFront, Glue), Azure, GCP, Kubernetes, Docker, Terraform, Jenkins, GitHub Actions, Serverless Framework

Data & Messaging: PostgreSQL, MySQL, MongoDB, DynamoDB, Apache Kafka, Apache Solr, Elasticsearch

AI/ML: LLM integration (Anthropic, OpenAI, Gemini), Model Context Protocol (MCP), prompt orchestration, PyTorch, Hugging Face, YOLOv8, collaborative filtering (ALS)

Frontend & Observability: React, Vue.js, AngularJS, Flutter, Grafana, Datadog, Kibana, Scalyr

PROFESSIONAL EXPERIENCE

Software Engineer Intern — Copart

May 2025 – Present · Dallas, TX

Fortune-500 online vehicle auction platform — Member Platform team

- **Cut server response times 60%** by restructuring API call patterns and applying Java design patterns across high-traffic Spring Boot services, deployed via Jenkins to a Kubernetes cluster
- **Built a company-wide AI integration framework** that lets teams add or edit LLM prompts through configuration or UI — no code changes or redeloys — accelerating AI feature rollout across the organization
- **Designed a high-throughput data pipeline** ingesting large volumes of member events into a datastore connected to Grafana, giving product teams real-time analytics
- **Reduced incident time-to-detection** by standing up in-house Apache Solr, Kafka, and message-queue services and integrating Scalyr, Datadog, and Grafana monitoring

Platform Engineer — Distinction Dev

Mar 2024 – Aug 2024

- **Cut resource consumption 40%** by building a real-time event fan-out platform with AWS Glue schema validation for a UK-based credit-card provider (fintech)
- **Reduced deployment failures 80%** by establishing end-to-end test pipelines (Vitest, Jest, Mocha) wired into GitHub Actions + Docker CI/CD
- **Eliminated duplicated integration code** by designing an Adapter-pattern layer (TypeScript, Inversify) that unified Salesforce, Mambu, and GoCardless behind one interface
- **Standardized legacy Lambda releases** by building an interim Infrastructure-as-Code framework with the AWS SDK and GitHub Actions, removing manual deployment steps

Software Developer — Distinction Dev

May 2023 – Mar 2024

- **Accelerated order processing 70%** by building serverless Node.js microservices on AWS (Serverless Framework, CloudFormation) for a UK retailer's nationwide supply-chain automation system tracking thousands of orders
- **Reduced SQL execution time 60%** by optimizing Sequelize ORM queries and MySQL indexing across order-tracking flows
- **Achieved ~40% faster page loads and 99.9% availability** by shipping Vue.js frontends with Auth0 authentication and an S3 / Route 53 / CloudFront delivery pipeline
- **Shortened release cycles by two weeks** by leading client demos, designing system architecture, and documenting interfaces and release processes

Flutter Developer Intern — UrbanWeb Solutions

Dec 2022 – May 2023

- **Delivered a cross-platform inventory-tracking app** (Flutter, PHP) giving a steel manufacturer's field staff real-time supply-chain visibility; led an intern team and ran client demos that secured next-phase approval

PROJECTS

Sentinel — AI-Powered Log Watchdog

github.com/dhyeyshah28801/sentinel

- Full-stack watchdog pulling infrastructure logs through a **custom MCP server** with pluggable adapters (Elasticsearch, DataSet/Scalyr, file); an **LLM classifies anomalies into structured JSON** and severity-gated alerts route to email/Slack — TypeScript monorepo, React dashboard, PostgreSQL/Drizzle

AI Orchestrator — Config-Driven LLM Framework

github.com/dhyeyshah28801

- Spring Boot (Java 21) framework where one endpoint executes any configured prompt against **Anthropic, OpenAI, or Gemini** via a provider-agnostic engine factory with Resilience4j retries and **JSON-Schema-validated output**; new prompts and provider swaps ship as config changes with **zero redeloys**

More projects — an explainable recommendation engine on 3.4M Instacart orders, a voice-to-SQL app, and a YOLOv8 parking-occupancy detector — at dhyeyshah28801.github.io/portfolio